

```
In [1]: from datascience import *
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plots
plots.style.use('fivethirtyeight')
```

```
In [2]: def draw_line(slope=0, intercept=0, x=None, color='r'):
    if x is None:
        x1, x2, y1, y2 = plots.gca().axis()
        x = make_array(x1, x2)
        y = x*slope + intercept
        plots.plot(x, y, color=color)
```

```
In [3]: def demographics_errors(slope, intercept):
    sample = [[14.7, 33995], [19.1, 61454], [50.7, 71183], [59.5, 105918]]
    demographics.scatter('College%', 'Median Income', alpha=0.5)
    xlims = make_array(5, 75)
    plots.plot(xlims, slope * xlims + intercept, lw=4)
    for x, y in sample:
        plots.plot([x, x], [y, slope * x + intercept], color='r', lw=4)
```

Regression Line

```
In [4]: def standard_units(arr):
    return (arr - np.average(arr))/np.std(arr)

def correlation(t, x, y):
    x_standard = standard_units(t.column(x))
    y_standard = standard_units(t.column(y))
    return np.average(x_standard * y_standard)

def slope(t, x, y):
    r = correlation(t, x, y)
    y_sd = np.std(t.column(y))
    x_sd = np.std(t.column(x))
    return r * y_sd / x_sd

def intercept(t, x, y):
    x_mean = np.mean(t.column(x))
```

```
y_mean = np.mean(t.column(y))
return y_mean - slope(t, x, y)*x_mean
```

```
In [5]: def fitted_values(t, x, y):
        """Return an array of the regressions estimates at all the x values"""
        a = slope(t, x, y)
        b = intercept(t, x, y)
        return a*t.column(x) + b
```

```
In [6]: demographics = Table.read_table('district_demographics2016.csv').drop(3)
        demographics.show(5)
```

State	District	Median Income	College%
Alabama	Congressional District 1 (115th Congress), Alabama	47083	24
Alabama	Congressional District 2 (115th Congress), Alabama	42035	21.8
Alabama	Congressional District 3 (115th Congress), Alabama	46544	22.8
Alabama	Congressional District 4 (115th Congress), Alabama	41110	17
Alabama	Congressional District 5 (115th Congress), Alabama	51690	30.3


... (430 rows omitted)

```
In [7]: demographics = demographics.drop(
        'State', 'District')
        demographics.show(5)
```

Median Income	College%
47083	24
42035	21.8
46544	22.8
41110	17
51690	30.3

... (430 rows omitted)

```
In [8]: demographics.scatter('College%', 'Median Income')
```



```
In [9]: correlation(demographics, 'College%', 'Median Income')
```

```
Out[9]: 0.8184648517141335
```


```
In [10]: regression_slope = slope(demographics, 'College%', 'Median Income')
         regression_intercept = intercept(demographics, 'College%', 'Median Income')
         (regression_slope, regression_intercept)
```

```
Out[10]: (1270.70168946388, 20802.577766677925)
```

```
In [11]: predicted = fitted_values(demographics, 'College%', 'Median Income')
```

```
In [12]: demographics = demographics.with_column(
```

```
'Linear Prediction', predicted)
demographics.scatter('College%')
```



```
In [13]: actual = demographics.column('Median Income')
          errors = actual - predicted
```

```
In [14]: demographics.with_column('Error', errors)
```

```
Out[14]:
```

Median Income	College%	Linear Prediction	Error
47083	24	51299.4	-4216.42
42035	21.8	48503.9	-6468.87
46544	22.8	49774.6	-3230.58
41110	17	42404.5	-1294.51
51690	30.3	59304.8	-7614.84
61413	36.7	67437.3	-6024.33


Median Income	College%	Linear Prediction	Error
34664	19.4	45454.2	-10790.2
76440	29.6	58415.3	18024.7
50537	24.5	51934.8	-1397.77
49072	34	64006.4	-14934.4

... (425 rows omitted)

```
In [15]: np.mean(errors ** 2) ** 0.5
```

```
Out[15]: 9398.515588571281
```

```
In [16]: demographics_errors(regression_slope, regression_intercept)
```




Root Mean Square Error

```
In [17]: def show_demographics_rmse(slope, intercept):  
    demographics_errors(slope, intercept)  
    x = demographics.column('College%')  
    y = demographics.column('Median Income')  
    prediction = slope * x + intercept  
    mse = np.mean((y - prediction) ** 2)  
    print("Root mean squared error:", mse ** 0.5)
```


```
In [18]: show_demographics_rmse(1500, 20000)
```

Root mean squared error: 11559.086490075999



```
In [19]: show_demographics_rmse(-1000, 75000)
```


Root mean squared error: 30247.883767944502



In [20]:


```
show_demographics_rmse(1500, 20000)
```

Root mean squared error: 11559.086490075999



```
In [21]: show_demographics_rmse(regression_slope, regression_intercept)
```


Root mean squared error: 9398.515588571281



Numerical Optimization

In [22]:

```
x = np.arange(1, 3, 0.1)
y = (x-2)**2 + 3
Table().with_columns('x', x, 'y', y).plot('x')
```



```
In [23]: def f(x):  
         return ((x-2)**2) + 3
```

```
In [24]: minimize(f)
```

```
Out[24]: 1.9999999946252267
```

Minimizing RMSE

```
In [25]: def demographics_rmse(any_slope, any_intercept):  
         x = demographics.column('College%')  
         y = demographics.column('Median Income')  
         estimate = any_slope*x + any_intercept  
         return (np.mean((y - estimate) ** 2)) ** 0.5
```

```
In [26]: minimize(demographics_rmse)
```

```
Out[26]: array([ 1270.70168805, 20802.57933807])
```

```
In [27]: make_array(regression_slope, regression_intercept)
```

```
Out[27]: array([ 1270.70168946, 20802.57776668])
```

Nonlinear Regression

```
In [28]: shotput = Table.read_table('shotput.csv')
```

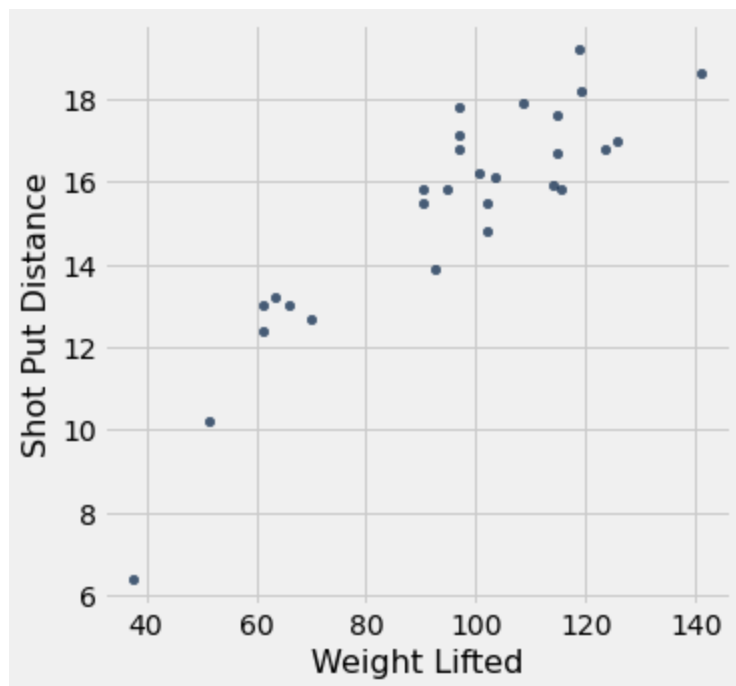
```
In [29]: shotput
```

```
Out[29]: Weight Lifted Shot Put Distance
```

37.5	6.4
51.5	10.2
61.3	12.4
61.3	13
63.6	13.2
66.1	13
70	12.7
92.7	13.9
90.5	15.5
90.5	15.8

... (18 rows omitted)

```
In [30]: shotput.scatter('Weight Lifted')
```




```
In [31]: def shotput_linear_rmse(any_slope, any_intercept):  
         x = shotput.column('Weight Lifted')  
         y = shotput.column('Shot Put Distance')  
         estimate = any_slope*x + any_intercept  
         return np.mean((y - estimate) ** 2) ** 0.5
```

```
In [32]: best_line = minimize(shotput_linear_rmse)  
         best_line
```

```
Out[32]: array([0.09834382, 5.95962883])
```

```
In [33]: weights = shotput.column(0)
```

```
In [34]: linear_fit = best_line.item(0)*weights + best_line.item(1)  
  
         shotput.with_column(  
             'Best Line', linear_fit  
         ).scatter(0)
```



Quadratic Function

$$f(x) = ax^2 + bx + c$$

for constants a , b , and c .

```
In [35]: def shotput_quadratic_rmse(a, b, c):
          x = shotput.column('Weight Lifted')
          y = shotput.column('Shot Put Distance')
          estimate = a*(x**2) + b*x + c
          return np.mean((y - estimate) ** 2) ** 0.5
```

```
In [36]: best_quad = minimize(shotput_quadratic_rmse)
          best_quad
```

```
Out[36]: array([-1.04003731e-03,  2.82706003e-01, -1.53167618e+00])
```


```
In [37]: # x = weight lifted = 100 kg
          # Then predicted shot put distance:
```

```
(-0.00104)*(100**2) + 0.2827*100 - 1.5318
```

Out[37]: 16.3382

```
In [38]: quad_fit = best_quad.item(0)*(weights**2) + best_quad.item(1)*weights + best_quad.item(2)
```

```
In [39]: shotput.with_column('Best Quadratic Curve', quad_fit).scatter(0)
```



In []:

In []: